

QBIN Specification (v1.0-draft)

Status: DRAFT

Version: 1.0 (file format major=1, minor=0)

Date: 2025-08-25

QBIN is a compact, binary container for OpenQASM circuits. It encodes gates and operands as typed binary records to minimize size and parsing overhead, while remaining extensible and round-trip safe to and from OpenQASM (2.x/3.x).

This document defines:

- File layout and section catalog
- Types, encodings, and alignment
- Instruction set and operand model
- Metadata and debugging facilities
- Versioning and compatibility
- Validation rules and error codes

Non-goals for v1 (can arrive in v1.1+ or v2):

- Full general-purpose program bytecode (loops, call stacks, heap)
- Arbitrary user-defined classical compute beyond simple compares

1. Terminology

- QASM: OpenQASM source text.
- QBIN: This binary format.
- Section: A typed region within the file.
- Varint: Unsigned LEB128 variable-length integer.
- ULEB128: Unsigned LEB128.
- SLEB128: Signed LEB128.
- BE/LE: Big-Endian / Little-Endian.
- Qubit index: Zero-based integer identifying a qubit.
- Bit index: Zero-based integer identifying a classical bit.

2. File overview

A QBIN file consists of:

Header (fixed-size)

Section Table (array of entries)

Sections... (payloads, each 8-byte aligned)

Sections are independent typed blobs, aligned to 8 bytes. This enables:

- Streaming decode of the instruction stream
- Lazy reading of metadata, strings, debug info
- Forward/backward compatibility via unknown-section skipping

3. Endianness and numeric conventions

- All fixed-width integers are Little-Endian (LE) unless a section explicitly overrides.
- Floating-point values in v1 are IEEE-754 binary32 (float32, LE).
- Angles are in radians.
- Complex numbers (if present in metadata) are encoded as two float32 values: real, imag.

4. File header

Offset	Size	Field
0x00	4	Magic = 0x51 0x42 0x49 0x4E (ASCII "QBIN")
0x04	1	Major version (0x01)
0x05	1	Minor version (0x00)
0x06	1	Flags:bit0 endianness (0=LE, 1=BE; v1 MUST be 0)bit1 section_table_hash_presentbits2..7
0x07	1	Header size (bytes). v1 = 24
0x08	4	Section count (u32, LE)
0x0C	4	Section table offset (u32, LE; bytes from file start)
0x10	4	Section table size (u32, LE)
0x14	4	Header CRC32 (u32, LE) over bytes [0x00..0x13]

Notes:

- Readers MUST verify magic and major_version.
- If flags.bit1 == 1, the Section Table trailer (see 5.3) includes a hash.

5. Section table

Each entry describes one section.

5.1 Section ID space

Section IDs are 32-bit tags. Recommended core set:

ID (u32)	ASCII	Meaning
0x4D455441	META	Metadata (key-value, binary)
0x53545253	STRS	String table (UTF-8)
0x51554253	QUBS	Qubit table
0x42495453	BITS	Classical bit table
0x50415253	PARS	Parameter table
0x47415445	GATE	Gate table
0x494E5354	INST	Instruction stream (mandatory)
0x44454247	DEBG	Debug info
0x5349474E	SIGN	Detached signature
0x43505253	CPRS	Compression dictionary (optional)
0x45585453	EXTS	Extensions registry

Vendor sections: use 0x56xxxxxx ("Vxxx") and namespace keys in META (e.g., "com.acme.*") to avoid collisions.

5.2 Section table entry

Offset	Size	Field
0x00	4	Section ID (u32)
0x04	4	Section offset (u32, LE; 8-byte aligned)
0x08	4	Section size (u32, LE)
0x0C	4	Section flags (u32): bit0 compressed, bit1 checksummed, others reserved

The table consists of `section_count` entries.

5.3 Section table trailer (optional)

Present only if `header.flags.bit1 == 1`:

Offset	Size	Field
+0x00	4	Hash algorithm (1=CRC32C, 2=xxh3_64)
+0x04	8	Hash value (u64; for CRC32C store in low 32 bits)

6. Types and encodings

- `u8/u16/u32/u64`: unsigned integers (LE).
- `varint`: ULEB128 (non-negative integers: indices, counts).
- `svarint`: SLEB128 (signed immediates).
- `f32`: IEEE-754 float32 (LE).

- bool: u8 (0 or 1).
 - Offsets: u32 LE from file start unless stated otherwise.
 - Strings (STRS): UTF-8, zero-terminated, deduplicated, referenced by varint string_id.
-

7. Core sections

7.1 STRS (String Table)

Purpose: deduplicate human-readable strings (names, metadata keys).

Layout:

```
u32 magic = 0x53545253 ("STRS")    // inside payload for self-check
u32 count
for i in [0..count):
    varint byte_len
    u8[byte_len] utf8_bytes
    u8 0x00 // terminator, not counted in byte_len
```

String IDs are implicit, 0..count-1. ID 0 SHOULD be the empty string "".

7.2 META (Metadata)

Purpose: small, structured metadata without requiring a JSON parser.

Encoding (compact KV map):

```
u32 magic = "META"
varint pair_count
repeat pair_count:
    varint key_str_id
    u8 value_type // 0=nil, 1=bool, 2=varint, 3=svarint, 4=f32, 5=str_id, 6=blob
    value as per type:
        bool: u8
        varint: ULEB128
        svarint: SLEB128
        f32: 4 bytes
        str_id: varint
        blob: varint len + u8[len]
```

Recommended keys (by convention, STRS entries): `qasm.version` ("2.0" or "3.0"), `source.name` (filename), `generator` (tool name/version), `target` (back-end name), `created.utc` (ISO-8601). Vendors SHOULD namespace keys (e.g., `com.acme.foo`).

7.3 QUBS (Qubit Table) [optional]

Purpose: map logical indices to named registers/ranges; optional geometry.

```
u32 magic = "QUBS"
varint qubit_count
u8 layout_present // 0/1
if layout_present:
    repeat qubit_count:
        f32 x, f32 y, f32 z // optional coordinates
varint alias_count
repeat alias_count:
    varint first_index
    varint count
    varint name_str_id // e.g., "q"
```

Without QUBS, qubits are [0..N) as referenced by INST.

7.4 BITS (Classical Bit Table) [optional]

```
u32 magic = "BITS"
varint bit_count
varint alias_count
repeat alias_count:
    varint first_index
    varint count
    varint name_str_id // e.g., "c"
```

7.5 PARS (Parameter Table) [optional]

Purpose: define symbolic parameters that may appear in instructions.

```
u32 magic = "PARS"
varint param_count
repeat:
    varint name_str_id // e.g., "theta"
    u8 kind // 0=angle(rad), 1=scalar, 2=duration(ns)
    u8 value_tag // 0=unbound, 1=const_f32, 2=expr_ref
    if value_tag==1: f32 value
    if value_tag==2: varint expr_id // future: expression DAG in EXTS
```

7.6 GATE (Gate Table) [optional]

Purpose: register custom gate signatures for decompile/validation.

```
u32 magic = "GATE"
varint decl_count
repeat:
    varint name_str_id // gate name
```

```

varint num_qubits
varint num_params
u8      flags           // bit0: opaque, bit1: unitary-known, etc.
varint body_len
u8[body_len] body_qbin  // optional inline QBIN sub-stream implementing gate

```

7.7 INST (Instruction Stream) [mandatory]

Purpose: the program body, a linear stream of encoded instructions.

7.7.1 Layout

```

u32  magic = "INST"
varint instr_count
repeat instr_count:
  u8  opcode
  u8  operand_mask      // bitmask indicates which operand slots follow
  ... operands          // present iff corresponding mask bit is 1

```

Operand slots (in order), presence indicated by operand_mask bits:

```

bit0: qubit_a      (varint)
bit1: qubit_b      (varint)
bit2: qubit_c      (varint)
bit3: angle_0      (tagged: f32 or param_ref)
bit4: angle_1      (tagged: f32 or param_ref)
bit5: angle_2      (tagged: f32 or param_ref)
bit6: param_ref    (varint param_id) // for gates with 1+ params or CALLG gate_id
bit7: aux_u32      (u32)           // misc (e.g., duration ns or bit index)

```

Angle slot encoding:

```

u8 tag    // 0 = literal f32, 1 = param_ref
if tag==0: f32 value
if tag==1: varint param_id

```

7.7.2 Opcodes (v1 core) Single-qubit:

```

0x01 X      qubit_a
0x02 Y      qubit_a
0x03 Z      qubit_a
0x04 H      qubit_a
0x05 S      qubit_a
0x06 SDG    qubit_a
0x07 T      qubit_a
0x08 TDG    qubit_a
0x09 SX     qubit_a
0x0A SXDG   qubit_a
0x0B RX     qubit_a, angle_0

```

```

0x0C RY      qubit_a, angle_0
0x0D RZ      qubit_a, angle_0
0x0E PHASE   qubit_a, angle_0
0x0F U       qubit_a, angle_0, angle_1, angle_2

```

Two-qubit and multi:

```

0x10 CX      qubit_a (control), qubit_b (target)
0x11 CZ      qubit_a, qubit_b
0x12 ECR     qubit_a, qubit_b
0x13 SWAP    qubit_a, qubit_b
0x14 CSX     qubit_a, qubit_b
0x15 CRX     qubit_a, qubit_b, angle_0
0x16 CRY     qubit_a, qubit_b, angle_0
0x17 CRZ     qubit_a, qubit_b, angle_0
0x18 CU      qubit_a, qubit_b, angle_0, angle_1, angle_2

0x20 RXX     qubit_a, qubit_b, angle_0
0x21 RYY     qubit_a, qubit_b, angle_0
0x22 RZZ     qubit_a, qubit_b, angle_0

```

IO and barriers:

```

0x30 MEASURE qubit_a, aux_u32(bit_index) // measure q -> c[bit_index]
0x31 RESET   qubit_a
0x32 BARRIER (no operands)

```

Timing and frame:

```

0x38 DELAY   qubit_a, aux_u32(duration_ns)
0x39 FRAME   qubit_a, angle_0 // frame change / phase ref

```

Custom gate call:

```

0x40 CALLG   param_ref(gate_id), qubit operands as per GATE entry
// operand_mask MUST include bit6 for gate_id, plus qubit bits (0..2) as needed

```

Reserved ranges:

- 0x80..0xBF: conditional and control flow (v1.1+)
- 0xC0..0xFF: vendor extensions (gated by EXTS)

7.7.3 Classical conditions (v1 minimal) A light-weight conditional wrapper encoded inline:

```

0x81 IF_EQ   aux_u32(bit_index), u8(value 0/1)
0x82 IF_NEQ  aux_u32(bit_index), u8(value 0/1)
0x8F ENDIF   (no operands)

```

Semantics:

- IF_* starts a guard that applies to subsequent instructions until ENDIF.

- Guards can nest; max nesting depth is implementation-defined but SHOULD be ≥ 8 .
- Maps to: `if (c[bit_index] == value) { ... }` in QASM3.

(Full while/for/switch are deferred to v1.1+ EXTS.)

8. Compression

If `section_flags.bit0` (compressed) is set, the section payload is wrapped:

```
u32 comp_magic = "CPRZ"
u8  alg        // 1=zstd, 2=lz4, 3=deflate
u32 raw_size
u8[] compressed_blob
```

Only the payload is compressed; the Section Table remains uncompressed.

9. Integrity and signatures

9.1 Section-level checksum (optional)

If a section entry has `section_flags.bit1` set, append:

```
u32 checksum_kind // 1=CRC32C
u32 checksum_value
```

Checksum covers the uncompressed payload bytes.

9.2 Detached signature (SIGN)

SIGN covers:

- header (0x00..0x13)
- section table bytes (including trailer if present)
- every section payload (uncompressed form)

Layout:

```
u32 magic = "SIGN"
u8  sig_kind // 1=CMS, 2=COSE_Sign1
varint sig_len
u8[sig_len] sig_bytes
```

10. Versioning and compatibility

- Readers **MUST** accept unknown sections by skipping via the Section Table.
- Readers **MUST** report unknown opcodes via `ERR_UNSUPPORTED_OPCODE` unless EXTS declares a decode hook.
- Minor versions (1.x) **MUST NOT** break existing encodings.
- Major version (2.0) may introduce breaking changes (distinct `major_version` value).

META **SHOULD** include:

```
qbin.version.major = 1
qbin.version.minor = 0
```

11. Validation rules

A file is valid if:

1. Magic, versions, and `header_crc32` are correct.
2. Section Table ranges are non-overlapping, 8-byte aligned, in-bounds.
3. INST section is present exactly once.
4. If QUBS/BITS are present, indices in INST are within declared counts.
5. If GATE is present and CALLG used, `gate_id` exists and operand counts match.
6. Angle and parameter references are well-typed.
7. Checksums (if present) validate.
8. Compression wrappers (if any) decode to declared `raw_size`.

Recommended additional checks:

- No NaN angles (or normalize to finite by policy).
 - Reasonable nesting depth for IF/ENDIF.
 - Optional max sizes (guard rails) to prevent resource abuse.
-

12. Error codes

Implementations **SHOULD** surface canonical errors:

```
0x01 ERR_MAGIC_OR_VERSION
0x02 ERR_HEADER_CRC
0x03 ERR_SECTION_TABLE_RANGE
0x04 ERR_MISSING_INST
0x05 ERR_MULTIPLE_INST
0x06 ERR_SECTION_CHECKSUM
0x07 ERR_DECOMPRESSION
```

```

0x08 ERR_TRUNCATED_SECTION
0x09 ERR_UNSUPPORTED_OPCODE
0x0A ERR_BAD_OPERAND_MASK
0x0B ERR_QUBIT_OOB
0x0C ERR_BIT_OOB
0x0D ERR_GATE_ID_OOB
0x0E ERR_PARAM_ID_OOB
0x0F ERR_GUARD_NESTING
0x10 ERR_TYPE_MISMATCH
0x11 ERR_META_FORMAT

```

13. Round-trip mapping to OpenQASM

13.1 QASM 3.0 example

QASM:

```

OPENQASM 3.0;
qubit[2] q;
bit[2] c;

h q[0];
cx q[0], q[1];
c[1] = measure q[1];
if (c[1] == 1) { x q[0]; }

```

QBIN INST (pseudo):

```

magic="INST"
count=6
1) H      a=0
2) CX     a=0, b=1
3) MEASURE a=1, aux(bit_index=1)
4) IF_EQ  aux(bit_index=1), value=1
5) X      a=0
6) ENDIF

```

13.2 Parameters

QASM:

```

cx q[0], q[1];
rz(pi/4) q[0];
rx(theta) q[1];

```

QBIN:

- RZ with angle_0 literal f32 = 0.785398163...

- RX with angle_0 param_ref to PARS entry "theta".
-

14. Debug info (DEBG)

Optional per-instruction mapping to original source:

```
u32  magic = "DEBG"
varint file_count
repeat file_count:
    varint path_str_id
varint mapping_count
repeat mapping_count:
    varint instr_index
    varint file_id
    varint line
    varint column
```

15. Extensions (EXTS) [for v1.1+]

If present, registers:

- Additional opcodes with operand schemas
- Expression DAGs for PARS (symbolic math)
- Vector operands (qubit lists) and multi-control gates

EXTS is a TLV catalog. Tools MAY ignore unknown tags.

16. Security considerations

- Treat all inputs as untrusted; cap sizes and counts.
 - Decompression bombs: enforce raw_size and compression ratio limits.
 - Signature verification: rely on standard CMS/COSE stacks.
 - Avoid executing embedded GATE bodies unless explicitly allowed.
-

17. Reference values

17.1 CRC32C (Castagnoli)

- Polynomial: 0x1EDC6F41
- Init: 0xFFFFFFFF
- XorOut: 0xFFFFFFFF

17.2 Angle normalization

- Writers SHOULD normalize angles to $[-\pi, \pi)$ to aid deduplication.
 - Readers MUST accept any finite float32.
-

18. Test vectors (informative)

Bell state minimal:

- META:
 - qasm.version = "3.0"
 - generator = "qbin-compiler 0.1"
- QUBS: qubit_count = 2
- INST:
 - 1) H a=0
 - 2) CX a=0, b=1

Expected sizes (indicative):

- QASM text: ~40-60 bytes
 - QBIN (uncompressed): ~24-48 bytes
 - QBIN (zstd): often ≤ 24 bytes
-

19. Conformance

A conforming encoder:

- Produces exactly one INST section.
- Emits a valid header and table.
- Ensures indices and counts are in range.
- Optionally emits STRS/META for better decompile fidelity.

A conforming decoder:

- Validates header and table.
 - Supports all v1 core opcodes and operand encodings.
 - Skips unknown sections and reports unknown opcodes.
-

20. License

This specification is provided under the MIT License, the same as the reference implementation in this repository.